



Heap Sort

Acknowledgement:

Donald E. Knuth, *The Art of Computer Programming, Volume 3: Sorting and Searching*, 2nd Ed., Addison Wesley, 1998, §5.2.3, p.144-8

Credit: Prof. Douglas Wilhelm Harder, ECE, University of Waterloo, Ontario, Canada


Outline

This topic covers the $\Theta(n \ln(n))$ sorting algorithm: *heap sort*

We will:


- define the strategy
- convert an unsorted list into a heap
- cover some examples
- analyze the run time complexity

Bonus: may be performed in place



Heap Sort

- Inserting n objects into a max-heap and then taking n objects will result in them coming out in order
- Strategy: Given an unsorted list with n objects, place them into a heap, and take them out



In-place Heapification

- Now, consider this unsorted array:

46	52	28	17	3	63	34	81	70	95
----	----	----	----	---	----	----	----	----	----


- Additionally, because arrays start at 0 (we started at entry 1 for binary heaps), we need different formulas for the children and parent

- The formulas are now:

Children of node k : Left child- $(2*k + 1)$,

Right child- $(2*k + 2)$


Parent of node k : $((k + 1)/2 - 1)$



Binary tree created out of array (**Not heap**)


In-place Heapification

Let's work bottom-up: each leaf node is a max heap on its own




In-place Heapification

- Starting at the back, we note that all leaf nodes are trivial heaps
- Also, the subtree with 87 as the root is a max-heap




In-place Heapification

- The subtree with 23 is not a max-heap, but swapping it with 55 creates a max-heap
- This process is termed *percolating down*.




In-place Heapification

The subtree with 3 as the root is not max-heap, but we can swap 3 and the maximum of its children: 86




In-place Heapification

Starting with the next higher level, the subtree with root 48 can be turned into a max-heap by swapping 48 and 99




In-place Heapification

Similarly, swapping 61 and 95 creates a max-heap of the next subtree




In-place Heapification

As does swapping 35 and 92




In-place Heapification

The subtree with root 24 may be converted into a max-heap by first swapping 24 and 86 and then swapping 24 and 28




In-place Heapification

The right-most subtree of the next higher level may be turned into a max-heap by swapping 77 and 99.




In-place Heapification

However, to turn the next subtree into a max-heap requires that 13 be percolated down to a leaf node.




In-place Heapification

The root need only be percolated down by two levels.



In-place Heapification

The final product is a max-heap.



Example Heap Sort


- Let us look at this example: we must convert the unordered array with $n = 10$ elements into a max-heap

46	52	28	17	3	63	34	81	70	95
----	----	----	----	---	----	----	----	----	----

- None of the leaf nodes need to be percolated down, and the first non-leaf node is in position $n/2$
- Thus we start with position $10/2 = 5$


Example Heap Sort

We compare 3 with its child and swap them




Example Heap Sort

We compare 17 with its two children and swap it with the maximum child (70)



Example Heap Sort


We compare 28 with its two children, 63 and 34, and swap it with the largest child



Example Heap Sort

We compare 52 with its children, swap it with the largest

- Recursing, no further swaps are needed



Example Heap Sort

Finally, we swap the root with its largest child, and recurse, swapping 46 again with 81, and then again with 70




Heap Sort Example

We have now converted the unsorted array

46	52	28	17	3	63	34	81	70	95
----	----	----	----	---	----	----	----	----	----

into a max-heap:




Heap Sort Example




Suppose we pop the maximum element of this heap

This leaves a gap at the back of the array:




Heap Sort Example



This is the last entry in the array, so why not fill it with the largest element?


Repeat this process: pop the maximum element, and then insert it at the end of the array:




Heap Sort Example

Repeat this process

➤ Pop and append 70




➤ Pop and append 63




Heap Sort Example

We have the 4 largest elements in order

➤ Pop and append 52




➤ Pop and append 46




Heap Sort Example

Continuing...

➤ Pop and append 34




➤ Pop and append 28



Heap Sort Example

Finally, we can pop 17, insert it into the 2nd location, and the resulting array is sorted





Black Board Example

Sort the following 12 entries using heap sort :

34, 15, 65, 59, 79, 42, 40, 80, 50, 61, 23, 46



Heap Sort Implementation

```
void heapSort(int arr[], int n)
{
    // Build heap (rearrange array)
    for (int i = n / 2 - 1; i >= 0; i--)
        heapify(arr, n, i);

    // One by one extract an element from heap
    for (int i=n-1; i>=0; i--)
    {
        // Move current root to end
        swap(arr[0], arr[i]);
        // call max heapify on the reduced heap
        heapify(arr, i, 0);
    }
}
```

Implementation

```
void heapify(int arr[], int n, int i)
{
    int largest = i;    // Initialize largest as root
    int l = 2*i + 1;    // left = 2*i + 1
    int r = 2*i + 2;    // right = 2*i + 2
    // If left child is larger than root
    if (l < n && arr[l] > arr[largest])
        largest = l;

    // If right child is larger than largest so far
    if (r < n && arr[r] > arr[largest])
        largest = r;

    if (largest != i) // If largest is not root
    {
        swap(arr[i], arr[largest]);
        // Recursively heapify the affected sub-tree
        heapify(arr, n, largest);
    }
}
```




Heap Sort

- Heapification runs in $\Theta(n)$
- Popping n items from a heap of size n , as we saw, runs in $\Theta(n \ln(n))$ time
 - We are only making one additional copy into the blank left at the end of the array
- Therefore, the total algorithm will run in $\Theta(n \ln(n))$ time

Run-time Analysis of Heapify

Considering a perfect tree of height h :

- The maximum number of swaps which a second-lowest level would experience is 1, the next higher level, 2, and so on



Run-time Analysis of Heapify

- At depth k , there are 2^k nodes and in the worst case, all of these nodes would have to percolated down $h - k$ levels
 - In the worst case, this would be requiring a total of $2^k(h - k)$ swaps
- Writing this sum mathematically, we get:

$$\sum_{k=0}^h 2^k (h - k) = (2^{h+1} - 1) - (h + 1)$$

Run-time Analysis of Heapify

- Recall that for a perfect tree, $n = 2^{h+1} - 1$ and $h + 1 = \lg(n + 1)$, therefore

$$\sum_{k=0}^h 2^k (h - k) = n - \lg(n + 1)$$

- Each swap requires two comparisons (which child is greatest), so there is a maximum of $2n$ (or $\Theta(n)$) comparisons

Run-time Analysis of Heapify

- Note that if we go the other way (treat the first entry as a max heap and then continually insert new elements into that heap, the run time is at worst

$$\begin{aligned}\sum_{k=0}^h 2^k k &= 2^{h+1} (h-1) + 2 \\ &= (2^{h+1} + 1)(h-1) - (h-1) + 2 \\ &= n(\lg(n+1) - 2) - \lg(n+1) + 4 = \Theta(n \ln(n))\end{aligned}$$

- It is significantly better to start at the back

Run-time Summary

The following table summarizes the run-times of heap sort

Case	Run Time	Comments
Worst	$\Theta(n \ln(n))$	No worst case
Average	$\Theta(n \ln(n))$	
Best	$\Theta(n)$	All or most entries are the same

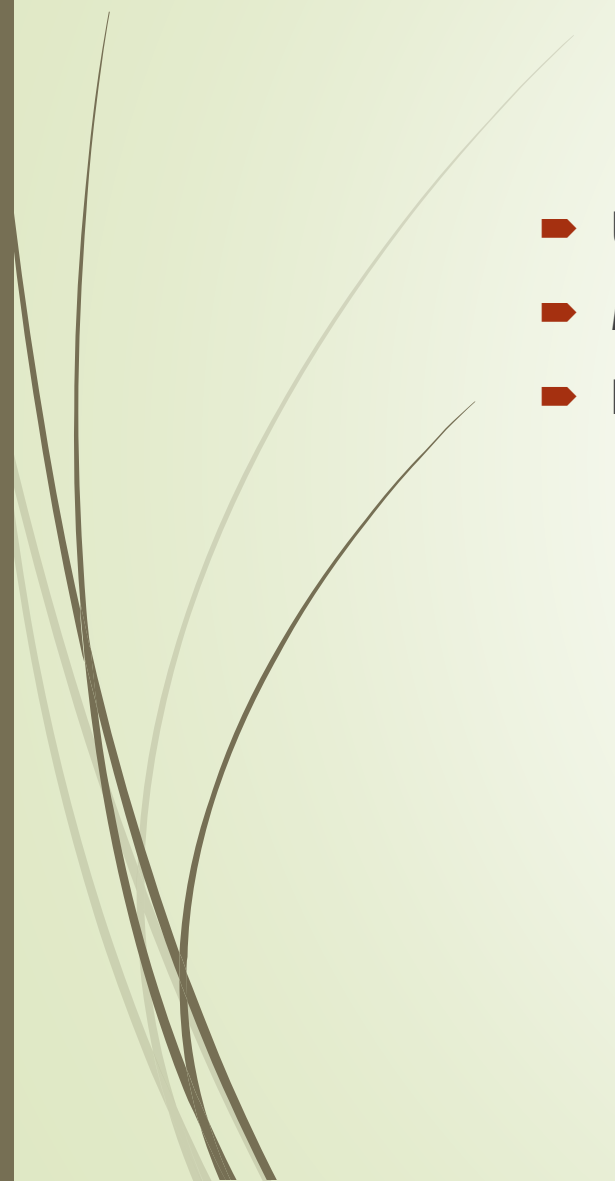



Use-cases

- Widely used in Database Management systems
 - implement B-Trees
 - sort SQL query results
- Operating Systems
 - Manage memory allocation – implementation of malloc and free functions
 - These functions dynamically allocate and deallocate memory in heap segment of OS
- Search engines
 - To sort query of results



Limitations

- Unstable sorting technique
 - Memory management is complex.
 - Has limited use in practice
- 



Summary

We have seen our first in-place $\Theta(n \ln(n))$ sorting algorithm:

- Convert the unsorted list into a max-heap as complete array
- Pop the top n times and place that object into the vacancy at the end
- It requires $\Theta(1)$ additional memory—it is truly in-place

It is a nice algorithm; however, we will see two other faster $n \ln(n)$ algorithms; however:

- Merge sort requires $\Theta(n)$ additional memory
- Quick sort requires $\Theta(\ln(n))$ additional memory