


Binary Search Trees

Binary Trees

◆ Recursive definition

1. An empty tree is a binary tree
2. A node with two child subtrees is a binary tree
3. Only what you get from 1 by a finite number of applications of 2 is a binary tree.

Is this a binary tree?




Binary Search Trees

- ♦ View today as data structures that can support **dynamic set operations**.
 - » Search, Minimum, Maximum, Predecessor, Successor, Insert, and Delete.
- ♦ Can be used to build
 - » **Dictionaries**.
 - » **Priority Queues**.
- ♦ Basic operations take time proportional to the height of the tree – **$O(h)$** .

BST – Representation


- ◆ Represented by a linked data structure of nodes.
- ◆ *root(T)* points to the root of tree T .
- ◆ Each node contains fields:
 - » *key*
 - » *left* – pointer to left child: root of left subtree.
 - » *right* – pointer to right child : root of right subtree.
 - » *p* – pointer to parent. $p[\text{root}[T]] = \text{NIL}$ (optional).



Binary Search Tree Property

- ◆ Stored keys must satisfy the *binary search tree* property.

- » $\forall y$ in left subtree of x , then $key[y] \leq key[x]$.
- » $\forall y$ in right subtree of x , then $key[y] \geq key[x]$.




Inorder Traversal

The binary-search-tree property allows the keys of a binary search tree to be printed, in (monotonically increasing) order, recursively.

Inorder-Tree-Walk (x)

1. **if** $x \neq \text{NIL}$
2. **then** Inorder-Tree-Walk($\text{left}[p]$)
3. print $\text{key}[x]$
4. Inorder-Tree-Walk($\text{right}[p]$)



- ◆ How long does the walk take?
- ◆ Can you prove its correctness?

Correctness of Inorder-Walk

- ◆ Must prove that it prints all elements, in order, and that it terminates.
- ◆ By induction on size of tree. Size=0: Easy.
- ◆ Size >1:
 - » Prints left subtree in order by induction.
 - » Prints root, which comes after all elements in left subtree (still in order).
 - » Prints right subtree in order (all elements come after root, so still in order).

Querying a Binary Search Tree

- ♦ All dynamic-set search operations can be supported in $O(h)$ time.
- ♦ $h = \Theta(\lg n)$ for a balanced binary tree (and for an average tree built by adding nodes in random order.)
- ♦ $h = \Theta(n)$ for an unbalanced tree that resembles a linear chain of n nodes in the worst case.


Tree Search

Tree-Search(x, k)

1. **if** $x = \text{NIL}$ *or* $k = \text{key}[x]$
2. **then** return x
3. **if** $k < \text{key}[x]$
4. **then** return Tree-Search($\text{left}[x], k$)
5. **else** return Tree-Search($\text{right}[x], k$)

Running time: $O(h)$


Aside: tail-recursion



Iterative Tree Search

Iterative-Tree-Search(x, k)

1. **while** $x \neq NIL$ **and** $k \neq key[x]$
2. **do if** $k < key[x]$
3. **then** $x \leftarrow left[x]$
4. **else** $x \leftarrow right[x]$
5. **return** x



The iterative tree search is more efficient on most computers.
The recursive tree search is more straightforward.

Finding Min & Max

- ♦ The binary-search-tree property guarantees that:
 - » The **minimum** is located at the **left-most** node.
 - » The **maximum** is located at the **right-most** node.

Tree-Minimum(x)

```
1. while  $left[x] \neq NIL$   
2.   do  $x \leftarrow left[x]$   
3. return  $x$ 
```

Tree-Maximum(x)

```
1. while  $right[x] \neq NIL$   
2.   do  $x \leftarrow right[x]$   
3. return  $x$ 
```

Q: How long do they take?

Predecessor and Successor

- ◆ Successor of node x is the **node y such that $key[y]$ is the smallest key greater than $key[x]$** .
- ◆ The successor of the largest key is NIL.
- ◆ Search consists of two cases.
 - » **If node x has a non-empty right subtree**, then x 's successor is the minimum in the right subtree of x .
 - » **If node x has an empty right subtree**, then:
 - As long as we move to the left up the tree (move up through right children), we are visiting smaller keys.
 - x 's successor y is the node that x is the predecessor of (x is the maximum in y 's left subtree).
 - In other words, x 's successor y , is the lowest ancestor of x whose left child is also an ancestor of x .


Pseudo-code for Successor

Tree-Successor(x)

- ♦ **if** $right[x] \neq NIL$
- 2. **then** return Tree-Minimum($right[x]$)
- 3. $y \leftarrow p[x]$
- 4. **while** $y \neq NIL$ **and** $x = right[y]$
- 5. **do** $x \leftarrow y$
- 6. $y \leftarrow p[y]$
- 7. **return** y


Code for *predecessor* is symmetric.

Running time: $O(h)$



BST Insertion – Pseudocode

- ◆ Change the dynamic set represented by a BST.
- ◆ Ensure the binary-search-tree property holds after change.
- ◆ Insertion is easier than deletion.



Tree-Insert(T, z)

1. $y \leftarrow \text{NIL}$
2. $x \leftarrow \text{root}[T]$
3. **while** $x \neq \text{NIL}$
4. **do** $y \leftarrow x$
5. **if** $\text{key}[z] < \text{key}[x]$
6. **then** $x \leftarrow \text{left}[x]$
7. **else** $x \leftarrow \text{right}[x]$
8. $p[z] \leftarrow y$
9. **if** $y = \text{NIL}$
10. **then** $\text{root}[t] \leftarrow z$
11. **else if** $\text{key}[z] < \text{key}[y]$
12. **then** $\text{left}[y] \leftarrow z$
13. **else** $\text{right}[y] \leftarrow z$

Analysis of Insertion

- ♦ Initialization: $O(1)$
 - ♦ While loop in lines 3-7 searches for place to insert z , maintaining parent y .
This takes $O(h)$ time.
 - ♦ Lines 8-13 insert the value: $O(1)$
- ⇒ TOTAL: $O(h)$ time to insert a node.

Tree-Insert(T, z)

```
1.   $y \leftarrow \text{NIL}$ 
2.   $x \leftarrow \text{root}[T]$ 
3.  while  $x \neq \text{NIL}$ 
4.      do  $y \leftarrow x$ 
5.          if  $\text{key}[z] < \text{key}[x]$ 
6.              then  $x \leftarrow \text{left}[x]$ 
7.              else  $x \leftarrow \text{right}[x]$ 
8.   $p[z] \leftarrow y$ 
9.  if  $y = \text{NIL}$ 
10.     then  $\text{root}[t] \leftarrow z$ 
11.     else if  $\text{key}[z] < \text{key}[y]$ 
12.         then  $\text{left}[y] \leftarrow z$ 
13.         else  $\text{right}[y] \leftarrow z$ 
```

Exercise: Sorting Using BSTs

Sort (A)

for $i \leftarrow 1$ to n

do tree-insert($A[i]$)

inorder-tree-walk($root$)

- » What are the worst case and best case running times?
- » In practice, how would this compare to other sorting algorithms?

Tree-Delete (T, x)

if x has no children ♦ case 0
 then remove x

if x has one child ♦ case 1
 then make $p[x]$ point to child

if x has two children (subtrees) ♦ case 2
 then swap x with its successor
 perform case 0 or case 1 to delete it

⇒ TOTAL: $O(h)$ time to delete a node

Deletion – Pseudocode

Tree-Delete(T, z)

/* Determine which node to splice out: either z or z 's successor. */

- ♦ **if** $left[z] = \text{NIL}$ **or** $right[z] = \text{NIL}$
- ♦ **then** $y \leftarrow z$
- ♦ **else** $y \leftarrow \text{Tree-Successor}[z]$

/* Set x to a non-NIL child of y , or to NIL if y has no children. */

4. **if** $left[y] \neq \text{NIL}$
5. **then** $x \leftarrow left[y]$
6. **else** $x \leftarrow right[y]$

/* y is removed from the tree by manipulating pointers of $p[y]$
and x */

7. **if** $x \neq \text{NIL}$
8. **then** $p[x] \leftarrow p[y]$

/* Continued on next slide */

Deletion – Pseudocode

Tree-Delete(T, z) (Contd. from previous slide)

```
9.   if  $p[y] = \text{NIL}$ 
10.      then  $\text{root}[T] \leftarrow x$ 
11.      else if  $y \leftarrow \text{left}[p[i]]$ 
12.          then  $\text{left}[p[y]] \leftarrow x$ 
13.          else  $\text{right}[p[y]] \leftarrow x$ 
/* If  $z$ 's successor was spliced out, copy its data into  $z$  */
14.  if  $y \neq z$ 
15.      then  $\text{key}[z] \leftarrow \text{key}[y]$ 
16.          copy  $y$ 's satellite data into  $z$ .
17.  return  $y$ 
```

Correctness of Tree-Delete

- ♦ How do we know case 2 should go to case 0 or case 1 instead of back to case 2?
 - » Because when x has 2 children, its successor is the minimum in its right subtree, and that successor has no left child (hence 0 or 1 child).
- ♦ Equivalently, we could swap with predecessor instead of successor. It might be good to alternate to avoid creating lopsided tree.

Binary Search Trees

- ♦ View today as data structures that can support **dynamic set operations**.
 - » Search, Minimum, Maximum, Predecessor, Successor, Insert, and Delete.
- ♦ Can be used to build
 - » **Dictionaries.**
 - » **Priority Queues.**
- ♦ Basic operations take time proportional to the height of the tree – **$O(h)$** .

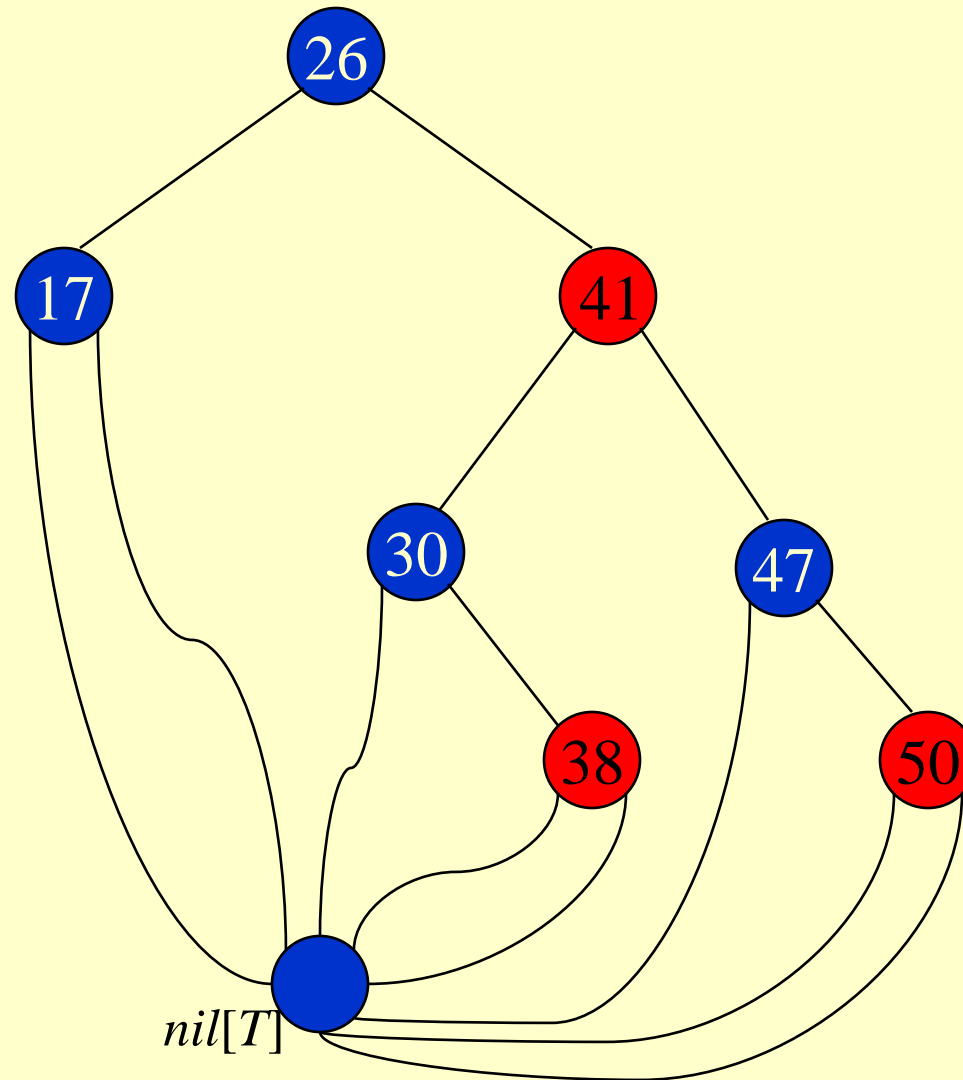
Red-black trees: Overview

- ♦ Red-black trees are a variation of binary search trees to ensure that the tree is *balanced*.
 - » Height is $O(\lg n)$, where n is the number of nodes.
- ♦ Operations take $O(\lg n)$ time in the *worst case*.

Red-black Tree

- ◆ Binary search tree + 1 bit per node: the attribute *color*, which is either **red** or **black**.
- ◆ All other attributes of BSTs are inherited:
 - » *key*, *left*, *right*, and *p*.
- ◆ All empty trees (leaves) are colored black.
 - » We use a single sentinel, *nil*, for all the leaves of red-black tree *T*, with *color*[*nil*] = black.
 - » The root's parent is also *nil*[*T*].

Red-black Tree – Example



Red-black Properties

1. Every node is either **red** or **black**.
2. The **root** is **black**.
3. Every **leaf** (*nil*) is **black**.
4. If a node is **red**, then both its children are **black**.
5. For each node, all paths from the node to descendant leaves contain the same number of **black** nodes.

Height of a Red-black Tree

- ◆ Height of a node:

- » Number of edges in a longest path to a leaf.

- ◆ Black-height of a node x , $bh(x)$:

- » $bh(x)$ is the number of black nodes (including $nil[T]$) on the path from x to leaf, not counting x .

- ◆ Black-height of a red-black tree is the black-height of its root.

- » By Property 5, black height is well defined.


Height of a Red-black Tree

♦ Example:

♦ Height of a node:

» Number of edges in a longest path to a leaf.

♦ Black-height of a node
 $bh(x)$ is the number of black nodes on path from x to leaf, not counting x .



Hysteresis : or the value of lazyness

- ♦ **Hysteresis**, n. [fr. Gr. to be behind, to lag.]
a retardation of an effect when the forces acting upon a body are changed (as if from viscosity or internal friction); *especially*: a lagging in the values of resulting magnetization in a magnetic material (as iron) due to a changing magnetizing force