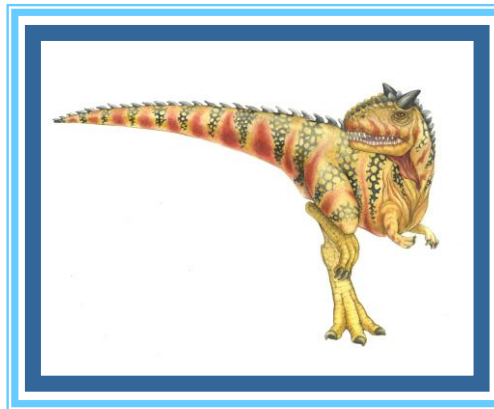


# Chapter 8: Virtual Memory

---





# Background

---

- ❑ Code needs to be in memory to execute, but entire program rarely used
  - ❑ Error code, unusual routines, large data structures
- ❑ Entire program code not needed at same time
- ❑ Consider ability to execute partially-loaded program
  - ❑ Program no longer constrained by limits of physical memory
  - ❑ Each program takes less memory while running -> more programs run at the same time
    - ▶ Increased CPU utilization and throughput
  - ❑ Less I/O needed to load or swap programs into memory -> each user program runs faster

**Virtual memory** – separation of user logical memory from physical memory





# Background (Cont.)

---

**Virtual address space** – logical view of how process is stored in memory

- Usually start at address 0, contiguous addresses until end of space
- Meanwhile, physical memory organized in page frames
- MMU must map logical to physical


Virtual memory can be implemented via:

- Demand paging
- Demand segmentation






# Virtual Memory That is Larger Than Physical Memory





# Demand Paging

- ❑ Could bring entire process into memory at load time
- ❑ Or bring a page into memory only when it is needed
  - ❑ Less I/O needed, no unnecessary I/O
  - ❑ Less memory needed
  - ❑ More users
- ❑ Page is needed  $\Rightarrow$  reference to it
  - ❑ not-in-memory  $\Rightarrow$  bring to memory
- ❑ **Lazy swapper** – never swaps a page into memory unless page will be needed
  - ❑ **pager** is concerned with the individual pages of a process





# Basic Concepts

---

- If pages needed are already **memory resident**
  - Execute the pages
- Otherwise, pager guesses which pages will be used before swapping out again
- Instead, pager brings in only those pages into memory
- If page needed and not memory resident
  - Need to detect and load the page into memory from storage
    - ▶ Without changing program behavior
    - ▶ Without programmer needing to change code





# Valid-Invalid Bit

- With each page table entry a valid–invalid bit is associated (**v**  $\Rightarrow$  in-memory – **memory resident**, **i**  $\Rightarrow$  not-in-memory)
- Initially valid–invalid bit is set to **i** on all entries
- Example of a page table snapshot:

Frame #	valid-invalid bit
	<b>v</b>
	<b>v</b>
	<b>v</b>
	<b>i</b>
...	
	<b>i</b>
	<b>i</b>

page table

- During MMU address translation, if valid–invalid bit in page table entry is **i**  $\Rightarrow$  **page fault**





# Page Table When Some Pages Are Not in Main Memory

0	A
1	B
2	C
3	D
4	E
5	F
6	G
7	H


logical  
memory

valid-invalid bit		
frame		
0	4	v
1		i
2	6	v
3		i
4		i
5	9	v
6		i
7		i

page table

0
1
2
3
4
5
6
7
8
9
10
11
12
13
14
15

physical memory







# Page Fault


---

- If Access to a page marked invalid causes a **page fault**
- Operating system looks at another table to decide:
  - Invalid reference  $\Rightarrow$  abort
  - Just not in memory
- 1. Find free frame
- 2. Swap page into frame via scheduled disk operation
- 3. Reset tables to indicate page now in memory  
Set validation bit = **v**
- 4. Restart the instruction that caused the page fault





# Steps in Handling a Page Fault





# Aspects of Demand Paging

- Extreme case – start process with *no* pages in memory
  - OS sets instruction pointer to first instruction of process, non-memory-resident -> page fault
  - And for every other process pages on first access
  - **Pure demand paging**
- Actually, a given instruction could access multiple pages -> multiple page faults
  - Consider fetch and decode of instruction which adds 2 numbers from memory and stores result back to memory
  - Pain decreased because of **locality of reference**
- Hardware support needed for demand paging
  - Page table with valid / invalid bit
  - Secondary memory (swap device with **swap space**)
  - Instruction restart





# Performance of Demand Paging (Cont.)

- Three major activities
  - Service the interrupt – careful coding means just several hundred instructions needed
  - Read the page – lots of time
  - Restart the process – again just a small amount of time
- Page Fault Rate  $0 \leq p \leq 1$ 
  - if  $p = 0$  no page faults
  - if  $p = 1$ , every reference is a fault
- Effective Access Time (EAT)

$$\begin{aligned} \text{EAT} = & (1 - p) \times \text{memory access} \\ & + p (\text{page fault overhead} \\ & \quad + \text{swap page out} \\ & \quad + \text{swap page in} ) \end{aligned}$$





# Demand Paging Example

- ❑ Memory access time = 200 nanoseconds
- ❑ Average page-fault service time = 8 milliseconds
- ❑  $EAT = (1 - p) \times 200 + p (8 \text{ milliseconds})$   
 $= 200 - p \times 200 + p \times 8,000,000$   
 $= 200 + p \times 7,999,800$
- ❑ If one access out of 1,000 causes a page fault, then  
EAT = 8200 nanoseconds.  
This is a slowdown by a factor of 40!!
- ❑ If want performance degradation < 10 percent
  - ❑  $220 > 200 + 7,999,800 \times p$   
 $20 > 7,999,800 \times p$
  - ❑  $p < .0000025$
  - ❑ < one page fault in every 400,000 memory accesses





# What Happens if There is no Free Frame?

---

- ❑ we assumed that each page faults at most once, when it is first referenced
- ❑ To increase our degree of multiprogramming, we are over-allocating memory
- ❑ Page fault → There are no free frames on the free-frame list
  - ❑ Algorithm – terminate? swap out? replace the page?
- ❑ Page replacement – find some page in memory, but not really in use, page it out
  - ❑ Performance – want an algorithm which will result in minimum number of page faults
- ❑ Same page may be brought into memory several times





# Basic Page Replacement

---


- Find the location of the desired page on disk
- Find a free frame
  1. If there is a free frame, use it
  2. If there is no free frame, use a page replacement algorithm to select a **victim frame**
  3. Write victim frame to disk if dirty
    - » Only modified pages are written to disk
- Bring the desired page into the (newly) free frame; update the page and frame tables
- Continue the process by restarting the instruction that caused the trap

Max 2 page transfers for page fault – increasing EAT





# Page Replacement







# Page Replacement Algorithms

## □ Page-replacement algorithm

- Want lowest page-fault rate on both first access and re-access
- Evaluate algorithm by running it on a particular string of memory references (reference string) and computing the number of page faults on that string
  - String is just page numbers, not full addresses
  - Repeated access to the same page does not cause a page fault
  - Results depend on number of frames available
- In all our examples, the **reference string** of referenced page numbers is

**7,0,1,2,0,3,0,4,2,3,0,3,0,3,2,1,2,0,1,7,0,1**





# First-In-First-Out (FIFO) Algorithm

- Reference string: **7,0,1,2,0,3,0,4,2,3,0,3,0,3,2,1,2,0,1,7,0,1**
- 3 frames (3 pages can be in memory at a time per process)

reference string

7 0 1 2 0 3 0 4 2 3 0 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2																		
	0	0	0																		
			1	1																	

page frames


15 page faults

- Can vary by reference string: consider 1,2,3,4,1,2,5,1,2,3,4,5
  - For four frames: 10 faults
  - For three frames: 9 faults
  - Adding more frames can cause more page faults!
    - ▶ **Belady's Anomaly**





# FIFO Illustrating Belady's Anomaly





# Optimal Algorithm

- Replace page that will not be used for longest period of time
  - 9 is optimal for the example
- How do you know this?
  - Can't read the future
- Used for measuring how well your algorithm performs

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2		2		2		2						7		
	0	0	0		0		4		0		0						0		
		1	1		3		3		3		1						1		

page frames





# Least Recently Used (LRU) Algorithm

- Use past knowledge rather than future
- Replace page that has not been used in the most amount of time
- Associate time of last use with each page

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2		4	4	4	0			1		1		1		
	0	0	0		0		0	0	3	3			3		0		0		
		1	1		3		3	2	2	2			2		2		7		

page frames

- 12 faults – better than FIFO but worse than OPT
- Generally good algorithm and frequently used





# LRU Algorithm (Cont.)

- Clock Counter implementation
  - Every page entry has a clock counter; every time page is referenced through this entry, copy the clock into the counter
  - When a page needs to be changed, look at the counters to find smallest value
    - ▶ Search through table needed
- Stack implementation
  - Keep a stack of page numbers in a doubly linked list form:
  - Page referenced:
    - ▶ move it to the top
    - ▶ Multiple pointers to be changed
  - But each update more expensive
- LRU and OPT are cases of **stack algorithms** that don't have Belady's Anomaly





# Counting Algorithms

---

- Keep a counter of the number of references that have been made to each page
  - Not common
- **Least Frequently Used (LFU) Algorithm**: replaces page with smallest count
- **Most Frequently Used (MFU) Algorithm**: based on the argument that the page with the smallest count was probably just brought in and has yet to be used





# Page-Buffering Algorithms

---

- Keep a pool of free frames, always
  - Then frame available when needed, not found at fault time
  - Read page into free frame and select victim to evict and add to free pool
  - When convenient, evict victim
- Possibly, keep list of modified pages
  - When backing store otherwise idle, write pages there and set to **non-dirty**







# Allocation of Frames

---

- Each process can be given ***minimum*** number of frames
- ***Maximum*** of course is total frames in the system
- Two major allocation schemes
  - fixed allocation
  - priority allocation
- Many variations





# Fixed Allocation

- Equal allocation – For example, if there are 100 frames (after allocating frames for the OS) and 5 processes, give each process 20 frames
  - Keep some as free frame buffer pool
- Proportional allocation – Allocate according to the size of process

- $s_i$  = size of process  $p_i$
- $S = \sum s_i$
- $m$  = total number of frames
- $a_i$  = allocation for  $p_i = \frac{s_i}{S} \times m$

$$m = 64$$

$$s_1 = 10$$

$$s_2 = 127$$

$$a_1 = \frac{10}{137} \cdot 62 \gg 4$$

$$a_2 = \frac{127}{137} \cdot 62 \gg 57$$





# Priority Allocation

---

- Use a proportional allocation scheme using priorities rather than size
- If process  $P_i$  generates a page fault,
  - select for replacement one of its frames
  - select for replacement a frame from a process with lower priority number





# Global vs. Local Allocation

---

- **Global replacement** – process selects a replacement frame from the set of all frames; one process can take a frame from another
  - But then process execution time can vary greatly
  - But greater throughput so more common
  
- **Local replacement** – each process selects from only its own set of allocated frames
  - More consistent per-process performance
  - But possibly underutilized memory





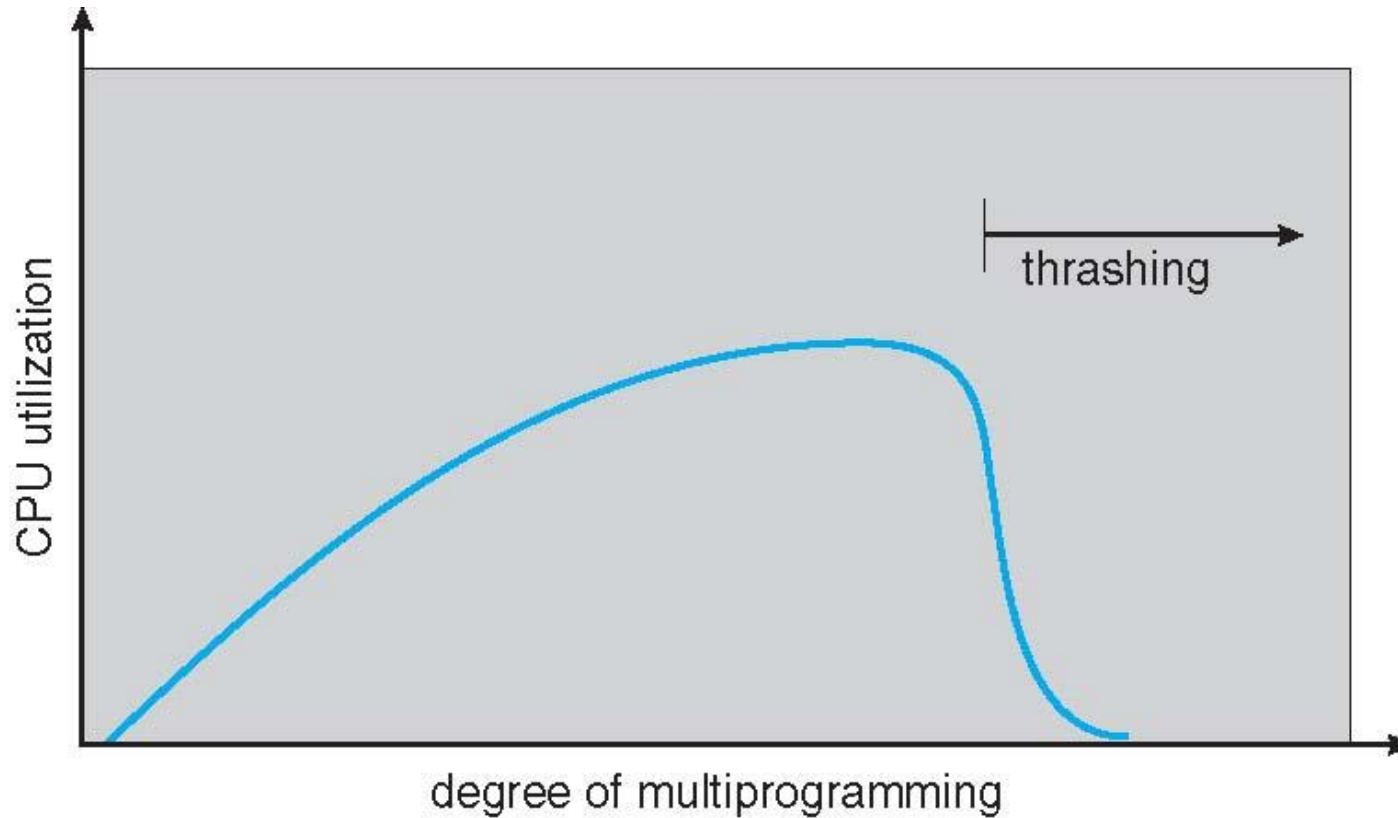
# Thrashing

- If a process does not have “enough” pages, the page-fault rate is very high
  - Page fault to get page
  - Replace existing frame
  - But quickly need replaced frame back
  - This leads to:
    - ▶ Low CPU utilization
    - ▶ Operating system thinking that it needs to increase the degree of multiprogramming
    - ▶ Another process added to the system
- **Thrashing**  $\equiv$  a process is busy swapping pages in and out





# Thrashing (Cont.)



# End of Chapter 8

---

