


Visual Attention

Visual Attention

- A model generated the caption “A woman is throwing a frisbee in a park” for the below image:



Visual Attention

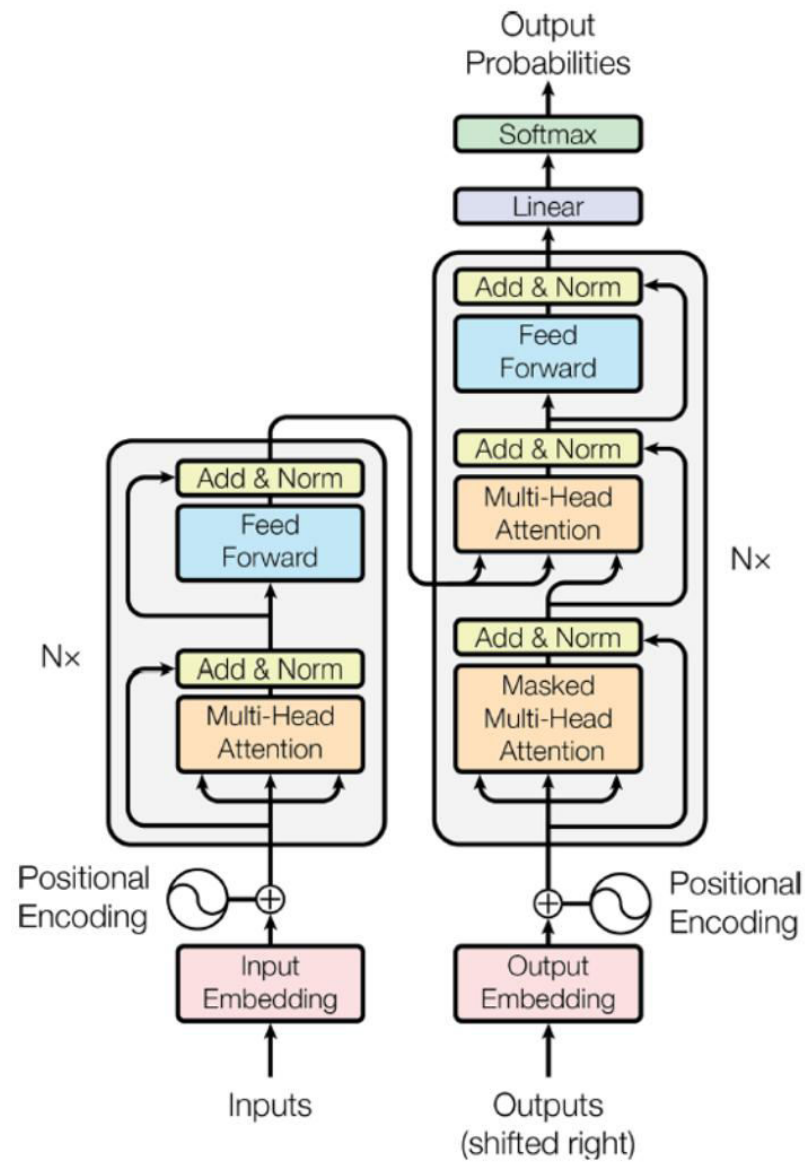
- CNN processes the image
- Outputs feature maps
- Then decoder RNN equipped with attention mechanism
 - Generates the caption, one word at a time

At each decoder time step (each word), the decoder uses the attention model to focus on just the right part of the image

The Transformer Architecture

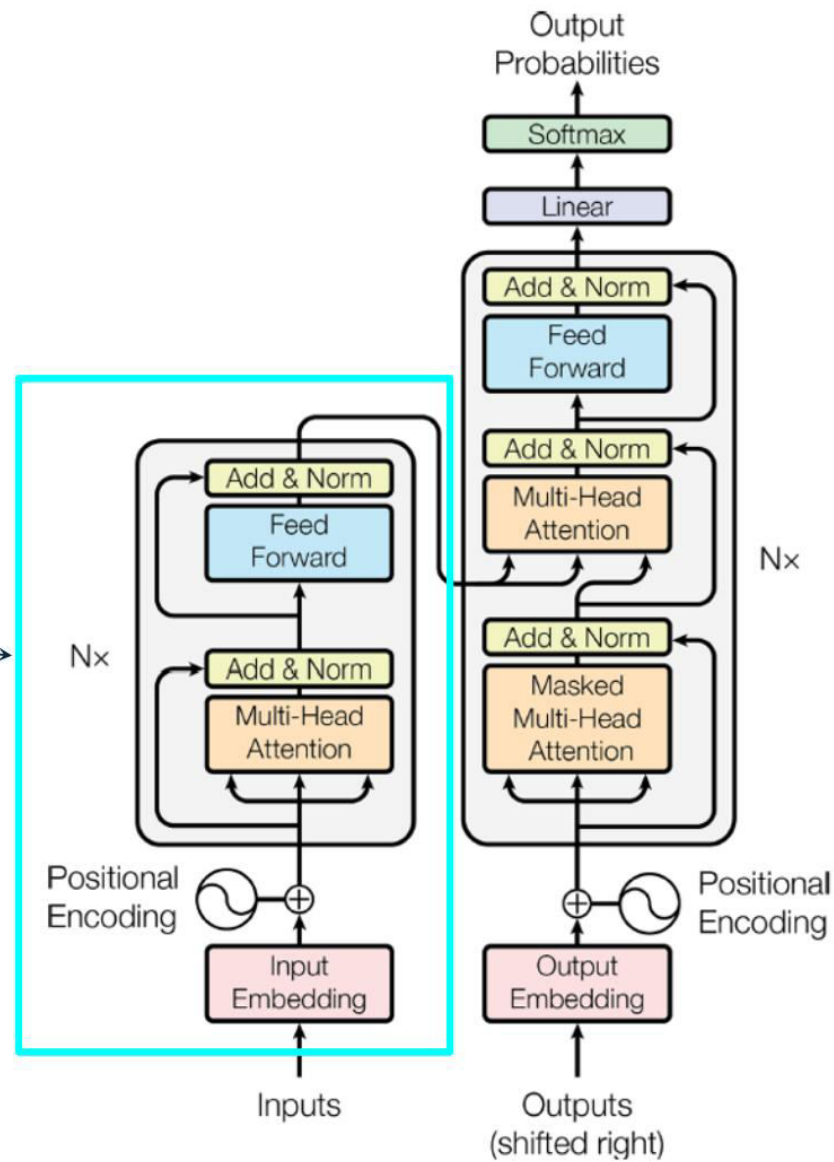
- In 2017, a team of Google researchers created an architecture called the **Transformer**
- It used only Attention Mechanism, no RNN or CNN to accompany it

The Transformer Architecture



The Transformer Architecture

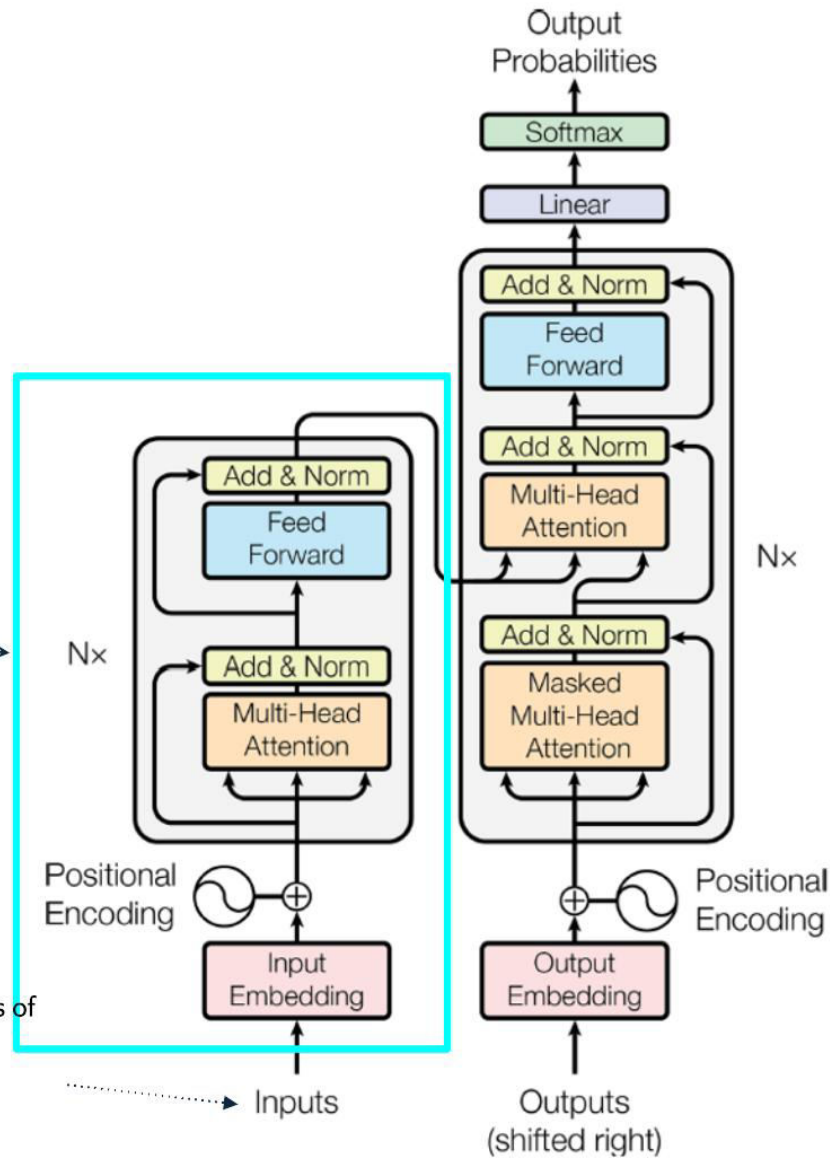
Encoder



The Transformer Architecture

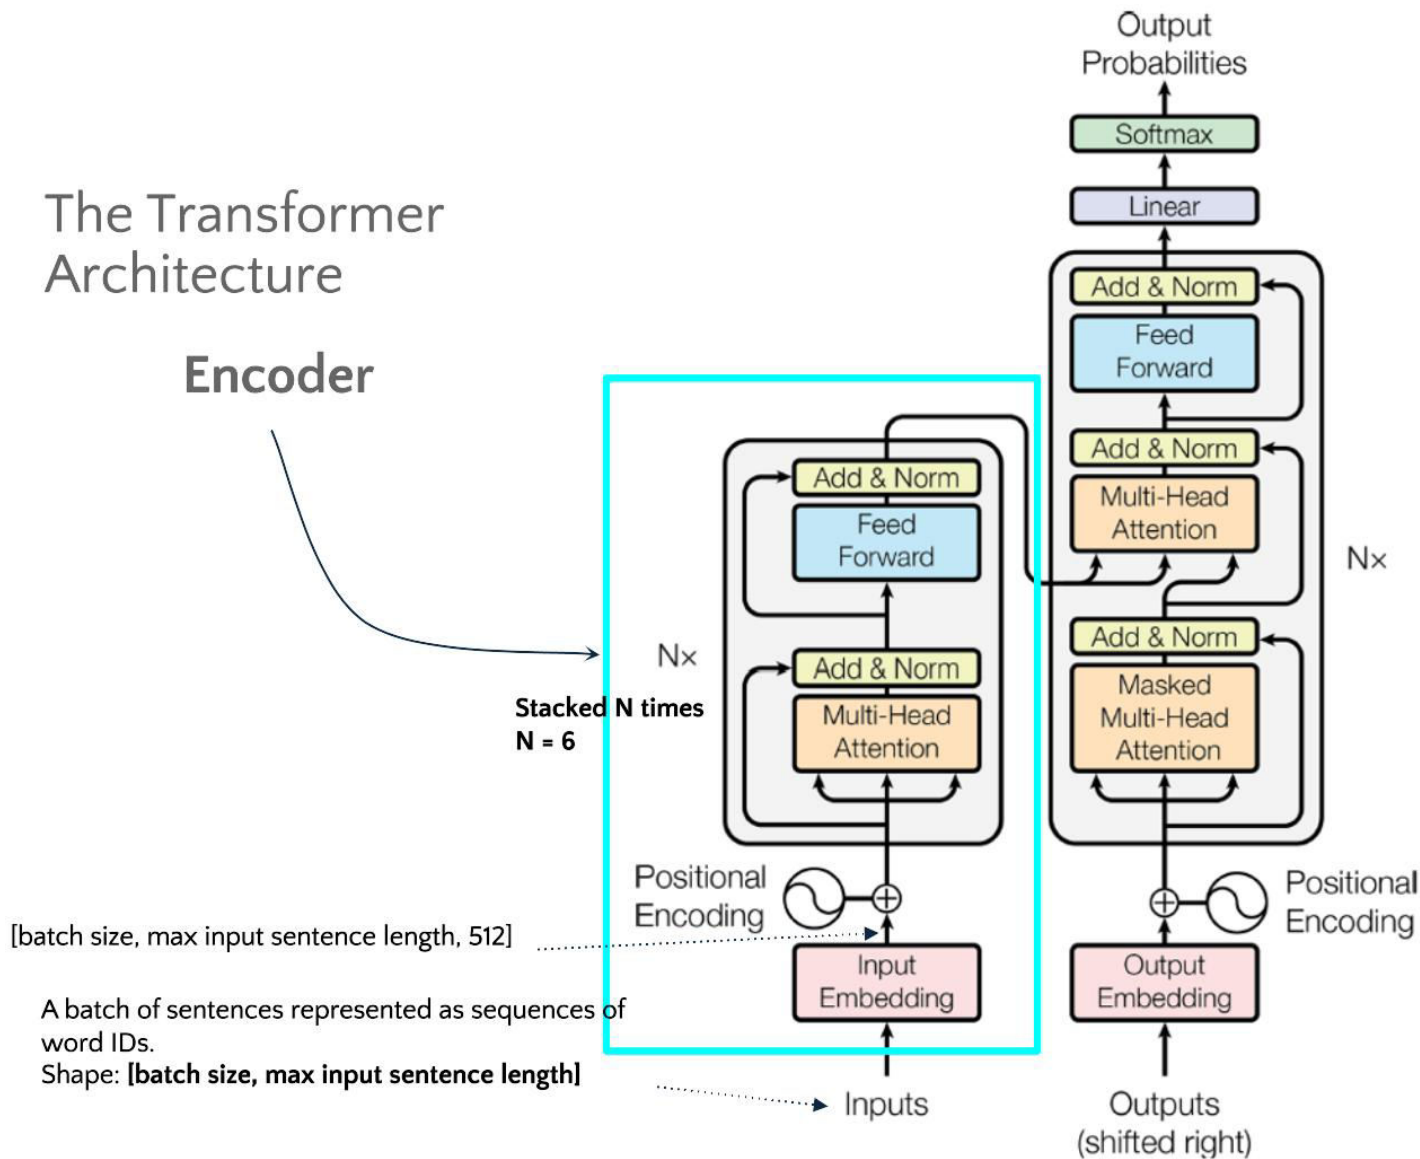
Encoder

A batch of sentences represented as sequences of word IDs.
Shape: [batch size, max input sentence length]

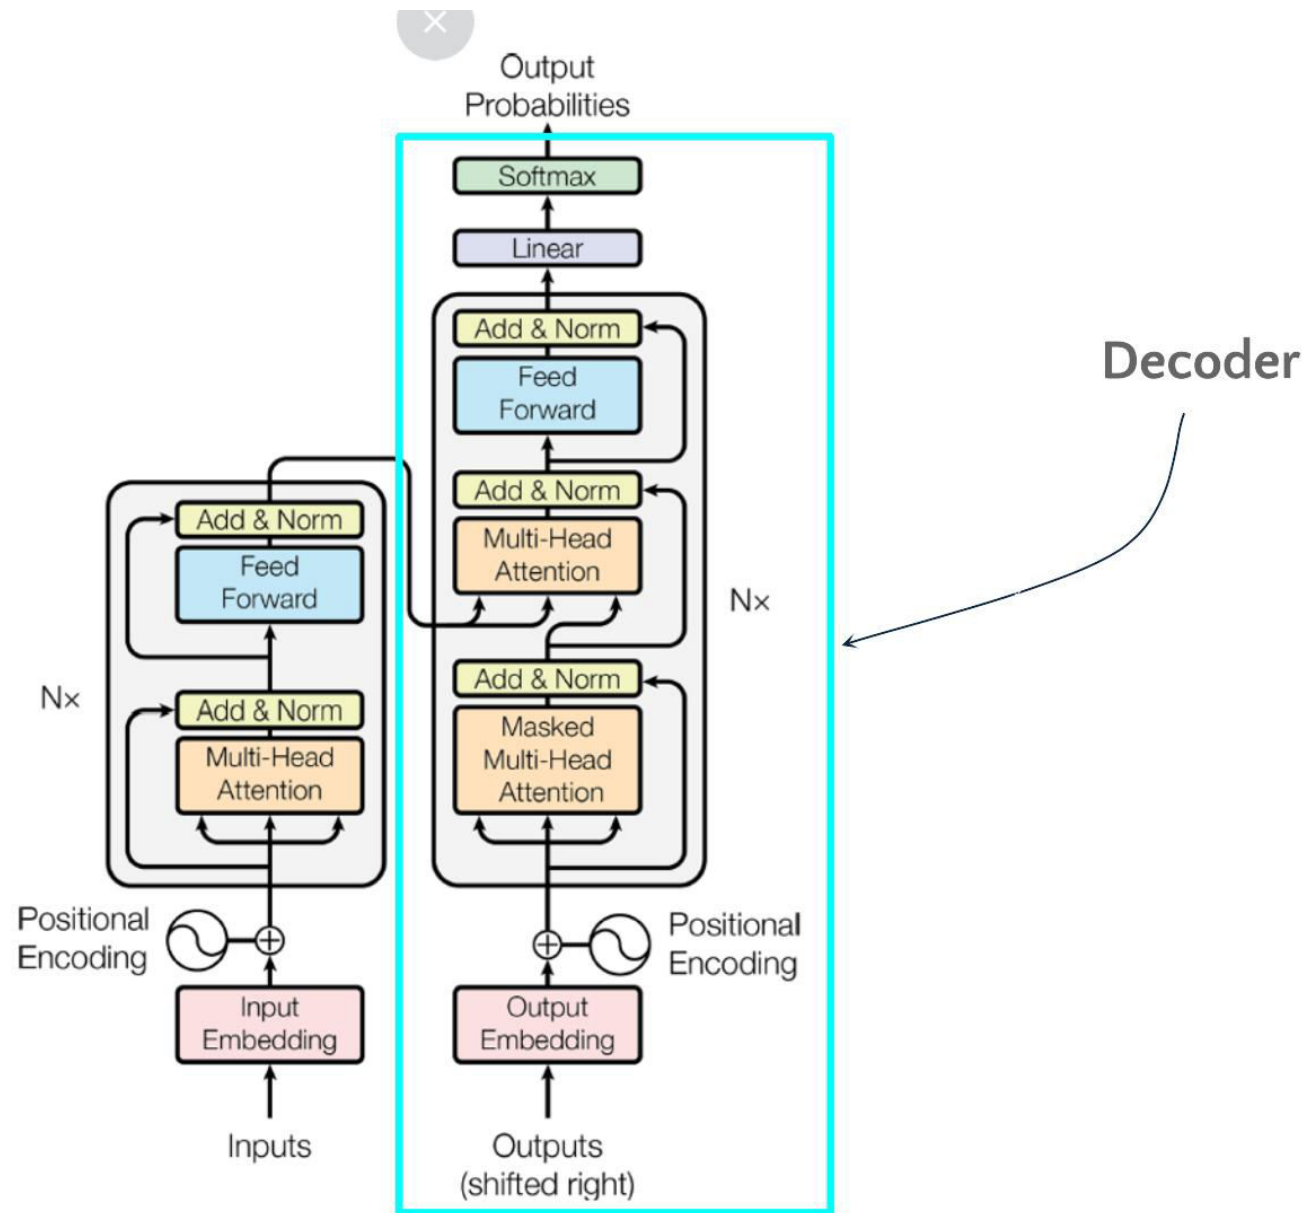


The Transformer Architecture

Encoder

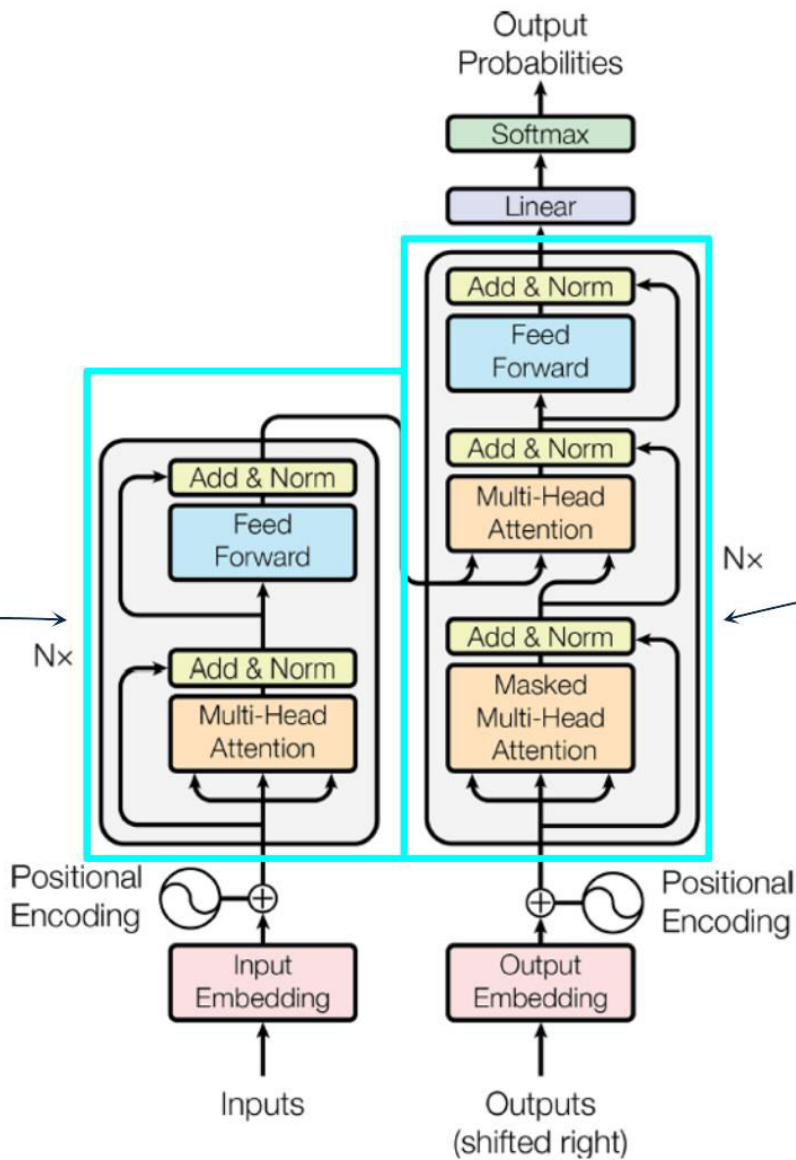


The Transformer Architecture



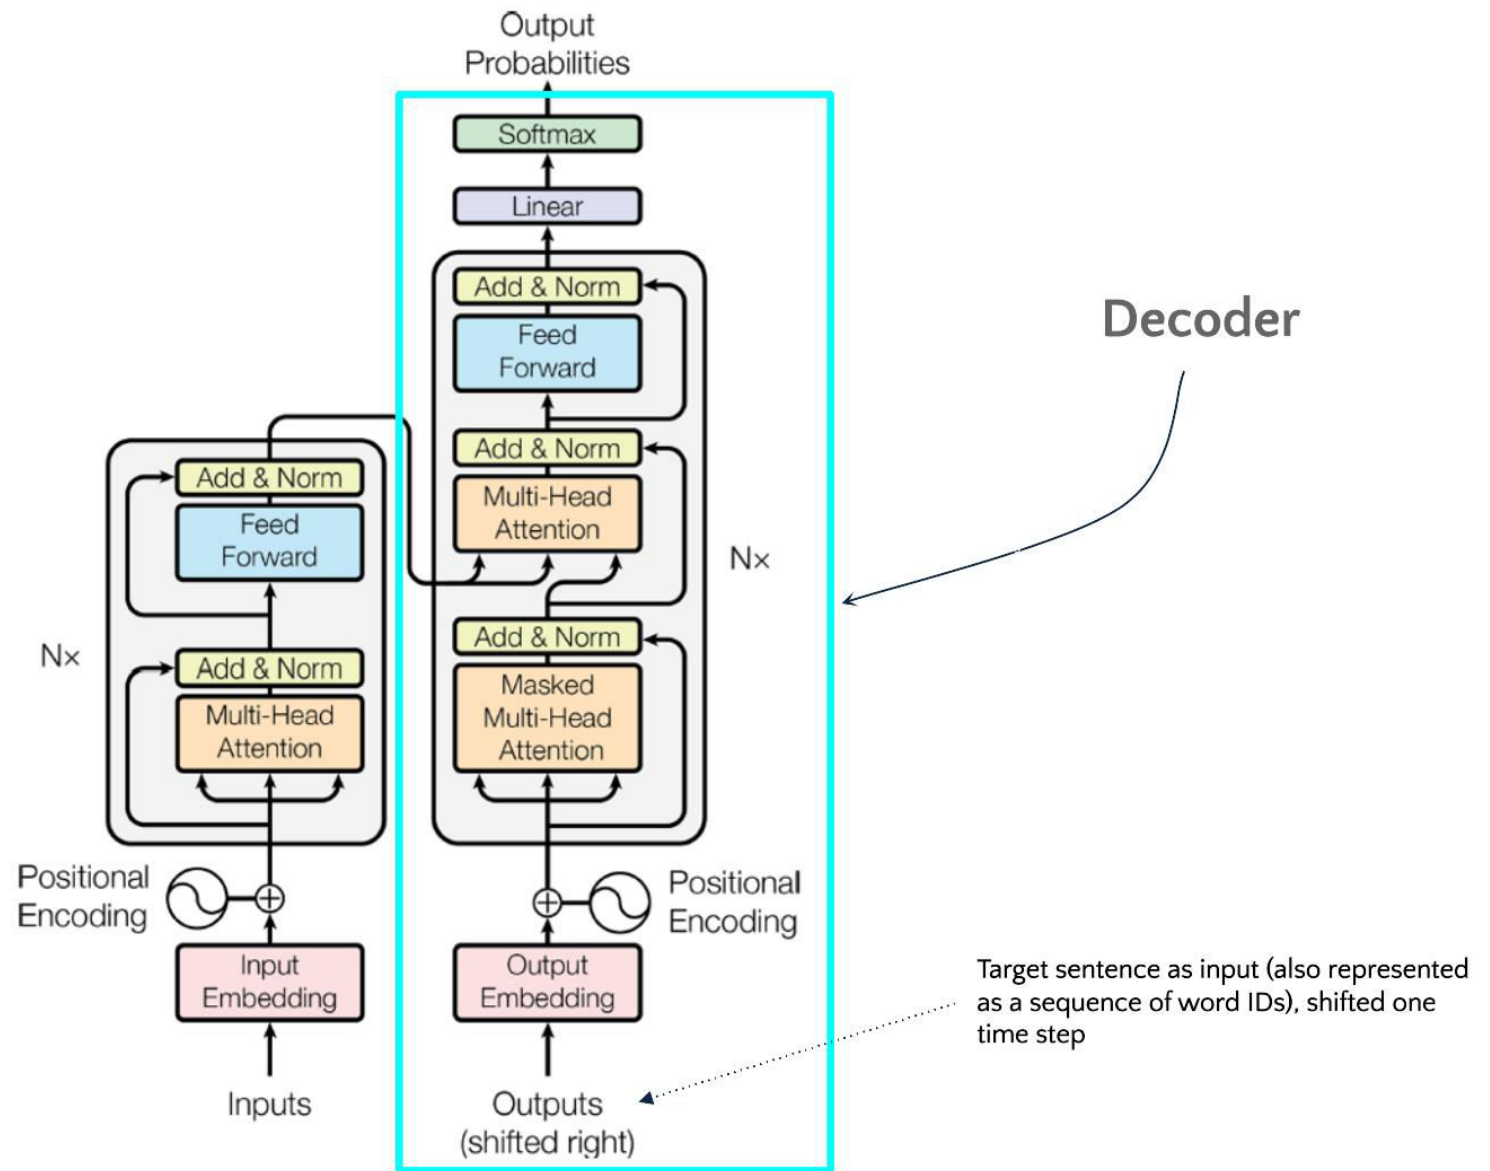
The Transformer Architecture

Encoder

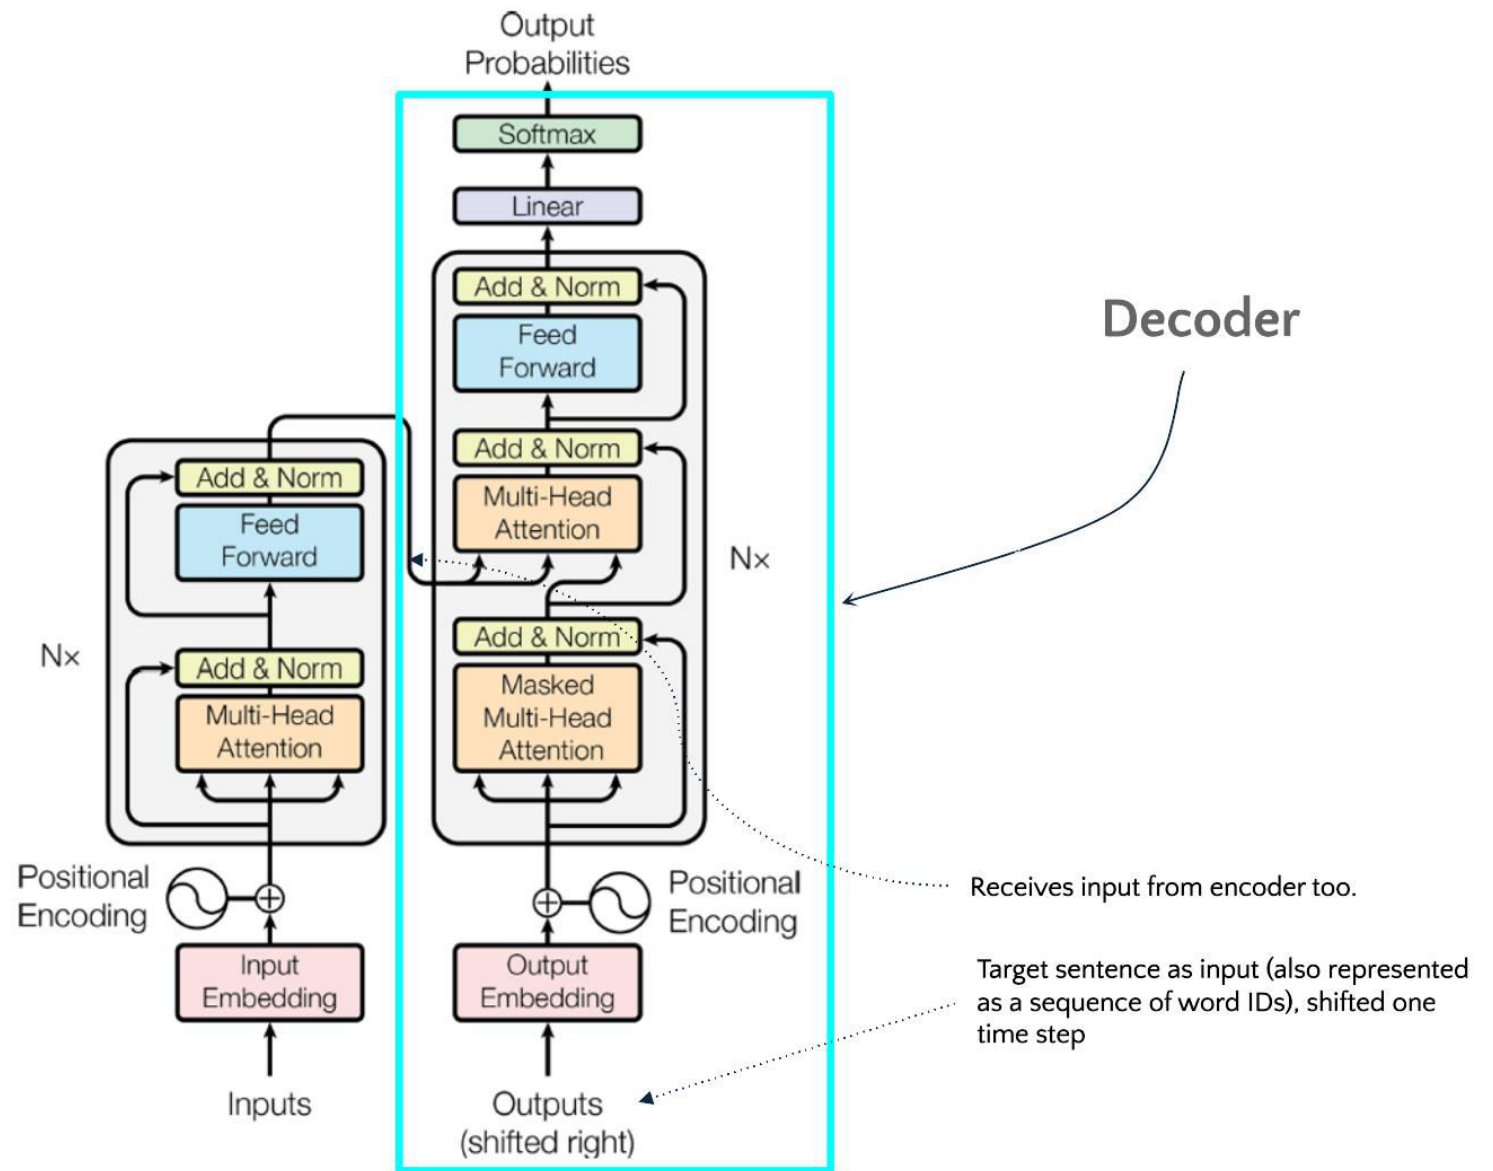


Decoder

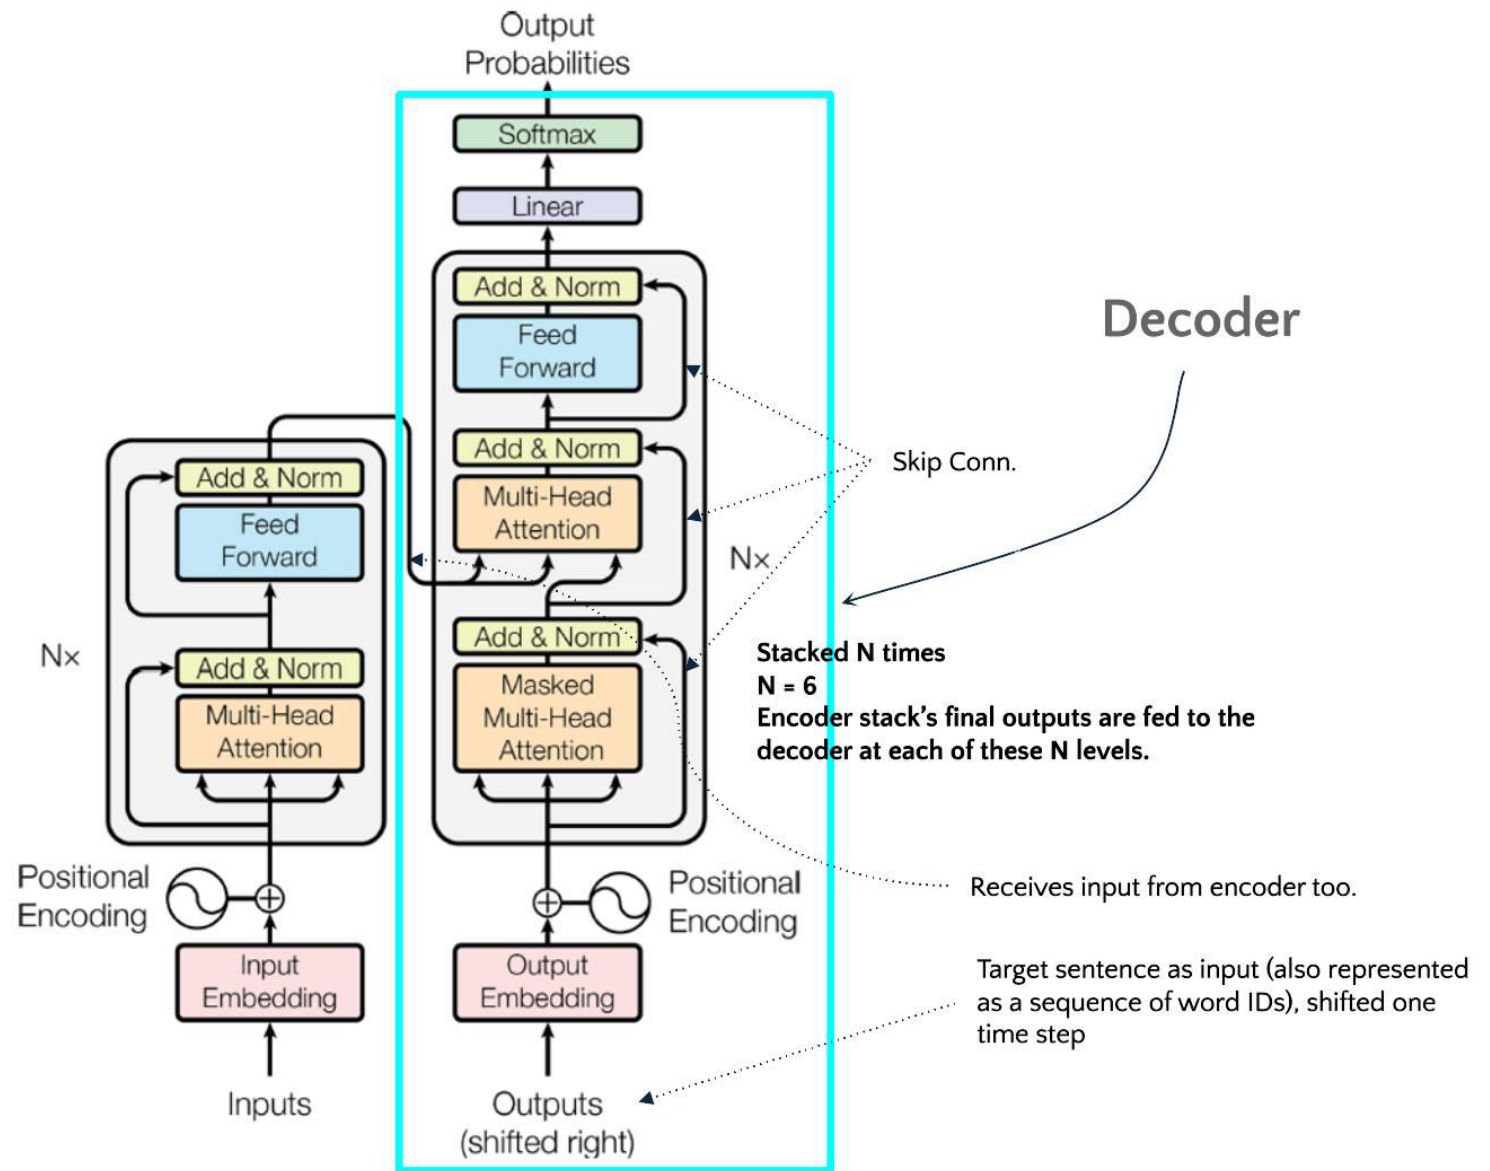
The Transformer Architecture



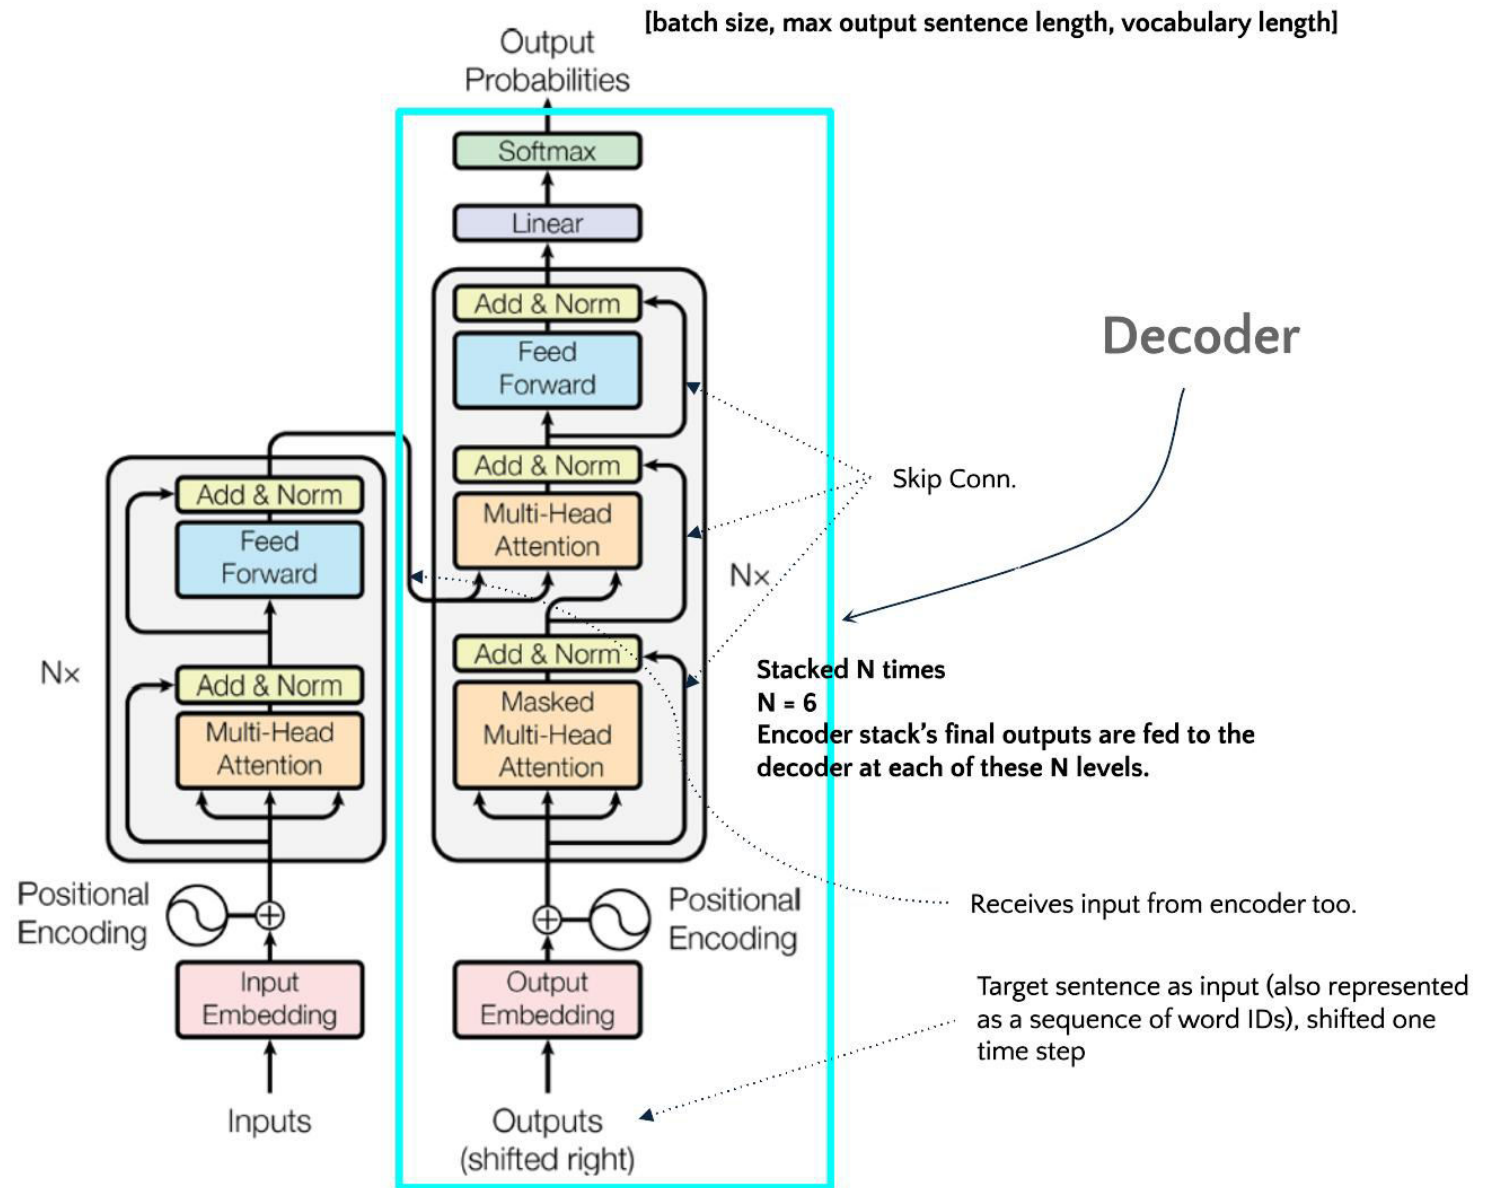
The Transformer Architecture



The Transformer Architecture

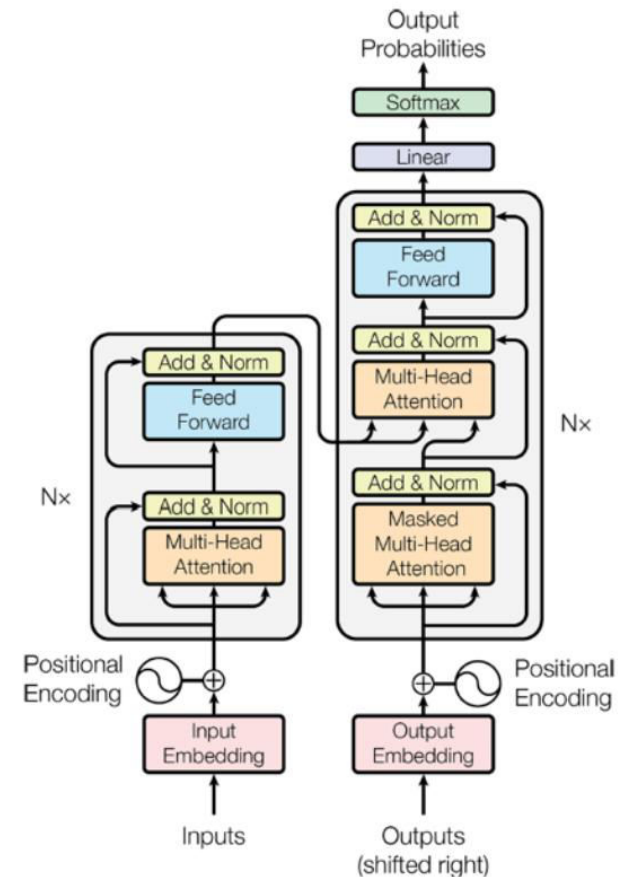


The Transformer Architecture



The Transformer Architecture (More)

- Two embedding layers
- $5 \times N$ skip connections
- Each of them followed by a layer normalization layer
- $2 \times N$ “Feed Forward” – 2 dense layers each
 - First one using the ReLU
 - second with no activation function
- The output layer is a dense layer using the softmax
- All of these layers are time-distributed.
 - So, each word is treated independently



The Transformer Architecture

- We are familiar with most components, except 2 of them:
 - Multi-Head Attention layer
 - Positional embeddings

Positional embeddings

References:

1. <https://arxiv.org/pdf/1706.03762.pdf>
2. https://kazemnejad.com/blog/transformer_architecture_positional_encoding/
3. https://www.tensorflow.org/tutorials/text/nmt_with_attention

Positional embeddings

Consider the 2 following sentences:

> I **do not** like the story of the movie, but I **do** like the cast

> I **do** like the story of the movie, but I **do not** like the cast

What is the difference between these 2 sentences?

Source: <https://is.gd/positionalembdding>

Multi-Head Attention

Multi-Head Attention

- Encodes each word's relationship with every other in a sentence
- Paying more attention to the most relevant ones
- Called **Self-Attention**

For example:

- “They welcomed the Queen of the United Kingdom”
- The output for “Queen” will depend on all the words in the sentence,
- but it will probably pay more attention to
 - “United” and “Kingdom” than
 - “They” or “welcomed.”

Multi-Head Attention

- Before we get into Multi-Headed Attention, we should first look at the concept of **Scaled Dot-Product Attention**
- **Example:**

Multi-Head Attention – Scaled Dot-Product Attention

- Say, encoder analyzed “They played chess,” and understood:
 - “They” → Subject
 - “Played” → verb
- This is encoded in the representations of the words
- Say, decoder has already translated the subject
 - and it thinks that it should translate the verb next
 - For this, it needs to fetch the verb from the input sentence.
- This is similar to a dictionary lookup:
 - Look up “verb” in {“subject”: “They”, “verb”: “played”, ...}

Multi-Head Attention

- Compared to the standard form of Attention, Scaled Dot-Product Attention utilizes Scaled Dot-Product to calculate similarity

$$\text{Attention}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{softmax}\left(\frac{\mathbf{Q}\mathbf{K}^T}{\sqrt{d_{\text{keys}}}}\right)\mathbf{V}$$

Multi-Head Attention

$$\text{Attention}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{softmax}\left(\frac{\mathbf{Q}\mathbf{K}^\top}{\sqrt{d_{\text{keys}}}}\right)\mathbf{V}$$

- \mathbf{Q} is a matrix containing one row per query of shape $[n_{\text{queries}}, d_{\text{keys}}]$
- n_{queries} is the number of queries
- d_{keys} is the number of dimensions of each query and each key

Multi-Head Attention

$$\text{Attention}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{softmax}\left(\frac{\mathbf{Q}\mathbf{K}^\top}{\sqrt{d_{\text{keys}}}}\right)\mathbf{V}$$

- \mathbf{K} is a matrix containing one row per key of shape $[n_{\text{keys}}, d_{\text{keys}}]$
- n_{keys} is the number of keys and values

Multi-Head Attention

$$\text{Attention}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{softmax}\left(\frac{\mathbf{Q}\mathbf{K}^\top}{\sqrt{d_{\text{keys}}}}\right)\mathbf{V}$$

- \mathbf{V} is a matrix containing one row per value of shape $[n_{\text{keys}}, d_{\text{values}}]$
- d_{values} is the number of each value

Multi-Head Attention

$$\text{Attention}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{softmax}\left(\frac{\mathbf{Q}\mathbf{K}^\top}{\sqrt{d_{\text{keys}}}}\right)\mathbf{V}$$

- The shape of $\mathbf{Q}\mathbf{K}^\top$ is $[n_{\text{queries}}, n_{\text{keys}}]$
- It contains one similarity score for each query/key pair

Multi-Head Attention

$$\text{Attention}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{softmax}\left(\frac{\mathbf{Q}\mathbf{K}^\top}{\sqrt{d_{\text{keys}}}}\right)\mathbf{V}$$

- The output of the softmax function has the same shape, but all rows sum up to 1
- The final output has a shape of $[n_{\text{queries}}, d_{\text{values}}]$, there is one row per query, where each row represents the query result (a weighted sum of the values)

Multi-Head Attention

- The `keras.layers.Attention` layer implements Scaled Dot-Product Attention
- Its inputs are just like Q, K, and V, except with an extra batch dimension (the first dimension)

Multi-Head Attention

- It is a bunch of Scaled Dot-Product Attention layers
- Each preceded by linear transformation of the values, keys, and queries
- All outputs are concatenated
- And they go through a final linear transformation

Multi-Head Attention

But why?

What is the intuition behind this architecture?

Multi-Head Attention

- The word representation encodes many different characteristics of the word
- With a single Scaled Dot-Product Attention layer this is not possible

Multi-Head Attention

- This is why Multi-Head Attention layer applies multiple different linear transformations of values, keys, and queries
- It gives the attention layer multiple “representation subspaces”



Next Word Prediction

Next Word Prediction

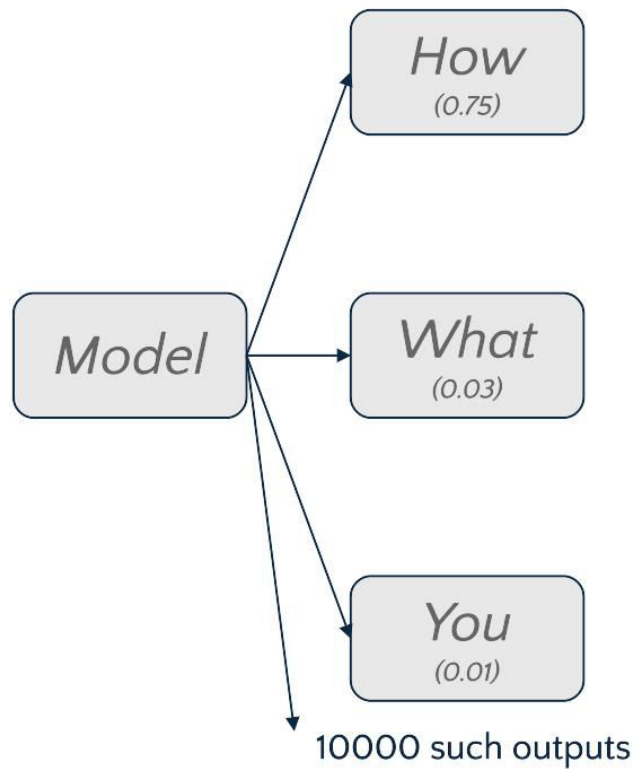
- Generating text requires
 - ensuring overall coherence,
 - grammar and
 - diversity in the sequence.
- A balance between local and global optimization.
 - Local optimization focuses on immediate choices,
 - while global optimization aims to find the overall best sequence.

Beam Search

Beam Search

- It keeps track of k (**beam width**) most promising sentences
- At each decoder step
 - It tries to extend them by one word
 - Keeping only the k most likely sentences

Beam Search (width - 3)




Beam Search: Working

Step 2

- Expansion:
 - Predict the next words to generate multiple possibilities.
 - Create new sequences by adding these words to the existing ones.

Beam Search (width - 3)



Beam Search: Working

Step 3

- Scoring:
 - Assign scores to sequences based on the probability of each word.
 - Calculate the overall probability for each sequence.